



北京大学  
PEKING UNIVERSITY

翰林81  
HIGH GO

# 密钥管理实践

开源开发实践-第六周

David & Cary

# CONTENT

## 目 录

- Introduction
- PostgreSQL Source Structure
- Adding a New Modules In PostgreSQL
- Internal KMS Considerations



```

    #deselection at the end - add back the deselected mirror modifier object
    mirror_ob.select= 1
    modifier_ob.select=1
    bpy.context.scene.objects.active = modifier_ob
    print("Selected" + str(modifier_ob)) # modifier ob is the active ob
    #mirror_ob.select = 0
    $one = bpy.context.selected_objects[0]
    #bpy.data.objects[one.name].select = 1
except:
    print("please select exactly two objects, the last one gets the modifier unless its not a mirror")

----- OPERATOR CLASSES -----
Mirror Tool

class MirrorX(bpy.types.Operator):
    """This adds an X mirror to the selected object"""
    bl_idname = "object.mirror_mirror_x"
    bl_label = "Mirror X"

    @classmethod
    def poll(cls, context):
        return context.active_object is not None
```




## Introduction



# Introduction

## What To Do Now?

- So far, we have learned the fundamentals and basics of a Key Management System (KMS).
- And the benefits of external KMS over internal.
- Now, it is time to start implementing an **internal KMS** as a separate module in PostgreSQL.
- We will walk through the Source structure of PostgreSQL and learn how you can add a new component to the system
- We will also discuss about some of the considerations in designing a KMS. What is it supposed to do?





```
        #deselection at the end - add back the deselected mirror modifier object
        mirror_ob.select= 1
        modifier_ob.select=1
        bpy.context.scene.objects.active = modifier_ob
        print("Selected" + str(modifier_ob)) # modifier ob is the active ob
        #mirror_ob.select = 0
        $one = bpy.context.selected_objects[0]
        #bpy.data.objects[one.name].select = 1
    except:
        print("please select exactly two objects, the last one gets the modifier unless its not a mod")
```

```
----- OPERATOR CLASSES -----
Mirror Tool
```

```
MirrorX(bpy.types.Operator):
    """This adds an X mirror to the selected object"""
    bl_idname = "object.mirror_mirror_x"
    bl_label = "Mirror X"
```

```
classmethod
def poll(cls, context):
    return context.active_object is not None
```



## PostgreSQL Source Structure



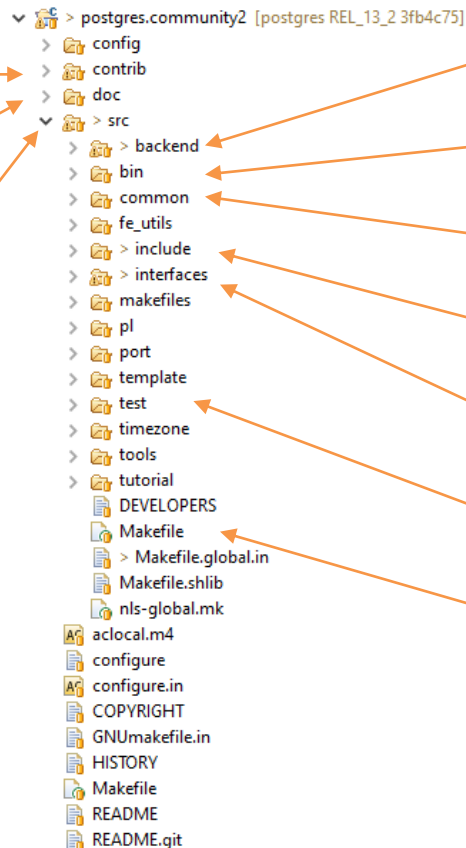
# PostgreSQL Root Source Structure

## The Root Source Structure

“**contrib**” folder contains the optional extensions that can be “added” onto PG. Will talk more in later class.

“**doc**” folder contains all the documentations of PG system written in **sgml**

“**src**” folder contains all the core components of PG. Both front end, backend, examples, tests, tools...etc



“**backend**” sub-folder contains all the backend components that drive PG.

“**bin**” sub-folder contains front end applications and tools that interacts with PG

“**common**” sub-folder contains common routines that everyone shares

“**include**” sub-folder contains all the header files for all the backend modules.

“**interfaces**” sub-folder contains interface-related or communication-related routines

“**test**” sub-folder contains all the regression test cases for PG. Will talk more in later class.

“**Makefile**” defines what to do after you issue the “make” command. Will discuss later



# PostgreSQL Frontend Source Structure



北京大学  
PEKING UNIVERSITY

清华  
HIGH GO

## The Frontend Source Structure

**“initdb”** folder contains initdb binary that initializes a new PG database cluster. This action is also called “bootstrapping”

**“pg\_basebackup”** folder contains pg\_basebackup binary that is responsible for taking a physical backup of your database

**“pg\_ctl”** folder contains pg\_ctl binary that allows you to start, stop, and restart a PG database

**“Makefile”** tells the “make” command to go into all the above folder and build each front end binaries

```
▼ postgres.community2 [postgres REL_13_2 3fb4c75]
> config
> contrib
> doc
▼ src
  > backend
  ▼ bin
    > initdb
    > pg_archivecleanup
    > pg_basebackup
    > pg_checksums
    > pg_config
    > pg_controldata
    > pg_ctl
    > pg_dump
    > pg_resetwal
    > pg_rewind
    > pg_test_fsync
    > pg_test_timing
    > pg_upgrade
    > pg_verifybackup
    > pg_waldump
    > pgbench
    > pgevent
    > psql
    > scripts
    > Makefile
```

**“pg\_controldata”** folder contains the pg\_controldata binary that allows you to check the current operating status of PG database

**“pg\_dump”** folder contains the pg\_dump binary that allows you to dump the data of the PG database.

**“pgbench”** folder contains the pgbench utility that allows you to simulate traffic load to test PG database performance

**“psql”** folder contains the psql utility that allows you to interact with a PG database

### Consideration:

Should KMS be added here as part of the front end?



PostgreSQL



# PostgreSQL Backend Source Structure



北京大学  
PEKING UNIVERSITY

清华大学  
HIGH GO

## The Backend Source Structure

**“access”** folder contains the access information for tuple data and indexes.  
Ex: heap, btree index, gist...etc

**“catalog”** folder contains the logics to interact with catalog tables

**“libpq”** folder contains the routines for communication between client and server

**“postmaster”** folder contains the main code for postmaster and how it manages child processes

**“Makefile”** tells the “make” command to compile and link all the backend components and produce the final output object

```
postgres.community2 [postgres REL_13_2 3fb4c75]
├── config
├── contrib
├── doc
└── src
    ├── backend
    │   ├── access
    │   ├── bootstrap
    │   ├── catalog
    │   ├── commands
    │   ├── executor
    │   ├── foreign
    │   ├── jit
    │   ├── lib
    │   ├── libpq
    │   ├── main
    │   ├── nodes
    │   ├── optimizer
    │   ├── parser
    │   ├── partitioning
    │   ├── po
    │   ├── port
    │   ├── postmaster
    │   ├── regex
    │   ├── replication
    │   ├── rewrite
    │   ├── snowball
    │   ├── statistics
    │   ├── storage
    │   ├── tcop
    │   ├── tsearch
    │   ├── utils
    │   ├── common.mk
    │   ├── Makefile
    │   └── nls.mk
    ├── bin
    ├── common
    └── fe_utils
```

**“bootstrap”** folder contains the initialization routines for PG. Such as during **“initdb”** process

**“executor”** folder contains logics to actually execute a **“query plan”**

**“parser”** folder contains logics to parse the SQL commands user enters.

**“replication”** folder contains logics to perform data replication between multiple PG instances

**“storage”** folder contains logics to handle data read and write between shared memory and the disk. This is the home of **“buffer manager”**

### Consideration:

Should KMS be added here as part of the backend components?



PostgreSQL



# The Makefile



北京大學  
PEKING UNIVERSITY

清華大學  
HIGH GO

## What To Do Now?

- You may have notice that there is a **Makefile** in both the **backend** and the **bin** folder and they behave very differently.
- In fact, there is almost always a Makefile in all folders and sub-folders under the **src directory**.
- So, what exactly is a Makefile?
- A **Makefile** is basically a series of instructions that tells the “**make**” command what to compile, how to compile, where to compile, how to install, which compiler to use, what flags to use...etc
- This file can get very, very complex! And one Makefile can include other Makefiles!
- This file is very important, especially if you want to add a new components inside PostgreSQL. You will have to have your own Makefile!
- Fortunately, you do not need to create a Makefile from scratch, you can copy the Makefile from another process and change it up a little!
- Before I talk a little more about Makefile, let me show you how compilation works

# The Compilation Process

## A Behind the Scene Look


That is the source file you will be working on. Normally ended with (.c) extension

Ex: bufmgr.c

Preprocessor then checks the syntax and compiler converts the expanded code into an assembly code

Then the assembly converts the assembly code into an object file. Normally every single (.c) file will produce one (.o) file.

Ex: bufmgr.o



Finally, the linker combines all of your object files (.o) including:

- Ones you compiled
- Ones you introduced to the system
- From shared and static libraries (.a) or (.so) files

Into one executable file. This is the final result of the compilation process.



# Makefile

## Makefile Basics

- The Makefile is a very big topic, and if you would like to learn more about Makefile, you can refer to this link:  
(<https://makefiletutorial.com/>)
- A Makefile consists of set of **targets** that you have to define what to do. A target name is set by the colon character (:)
- After the (:), is the prerequisites, or requirements for this target. Normally the command is a series of Bash scripts and compiler commands

```
targets: prerequisites  
command  
command  
command
```

- Common variables that we use in Makefiles

Variable	Meaning
CC	The compiler to use (ex: gcc)
CFLAGS	The compiler flags to use (ex: -g -O0)
OBJS	List of object files required to build (ex: bufmgr.o buf_init.o freelist.o)
LDFLAGS	Additional paths to locate additional libraries during linking process. (ex: -L/usr/local/lib/ -L/usr/lib)
LIBS	Actual libraries that we want to link. (ex: -lcrypto -lssl )
INCLUDE	Additional paths to locate the header files during compilation process. (ex: -I/usr/local/include -I/usr/include)

# Makefile Example

## Basic Makefile Example

```
all: hg_smgrserver hg_smgrserver_test

hg_smgrserver: hg_smgrserver.o hg_netio_utils.o
    $(CC) $(CFLAGS) $(OBJS) $(LDFLAGS) $(LDFLAGS_EX) $(LIBS) -o hg_smgrserver

hg_smgrserver_test: hg_smgrclient_test.o
    $(CC) $(CFLAGS) hg_smgrclient_test.o $(LDFLAGS) $(LDFLAGS_EX) $(LIBS) -o hg_smgrclient_test

install: all
    $(INSTALL_PROGRAM) hg_smgrserver$(X) '$(DESTDIR)$(bindir)'/hg_smgrserver$(X)

clean distclean maintainer-clean:
    rm -f hg_smgrserver$(X) $(OBJS) hg_smgrclient_test$(X) hg_smgrclient_test.o
```

For example:

- The command “**make**” will trigger the “**all**” target to be executed. This is the default target since we do not specify a target
- “**all**” requires 2 other targets to be done, **hg\_smgrserver** and **hg\_smgrserver\_test**, so they will get run as well if they haven’t been created
- **hg\_smgrserver** requires 2 object files and **hg\_smgrserver\_test** requires 1 object file. If the dependencies are met the build commands will be run to produce the final executable.



For example:

- The command “**make install**” will trigger the “**install**” target to be executed, which requires “**all**” target to be completed already. Then it will run the commands to install
- The command “**make clean**” will trigger the “**clean**” target to be executed, which does not depend on anything and it will just run the commands to clean up the code.



```
        #deselection at the end - add back the deselected mirror modifier object
        mirror_ob.select = 1
        modifier_ob.select = 1
        bpy.context.scene.objects.active = modifier_ob
        print("Selected" + str(modifier_ob)) # modifier ob is the active ob
        #mirror_ob.select = 0
        #one = bpy.context.selected_objects[0]
        #bpy.data.objects[one.name].select = 1
    except:
        print("please select exactly two objects, the last one gets the modifier unless its not a mod")
```

----- OPERATOR CLASS -----  
Mirror Tool

```
class MirrorX(bpy.types.Operator):
    """This adds an X mirror to the selected object"""
    bl_idname = "object.mirror_mirror_x"
    bl_label = "Mirror X"
```

```
    @classmethod
    def poll(cls, context):
        return context.active_object is not None
```

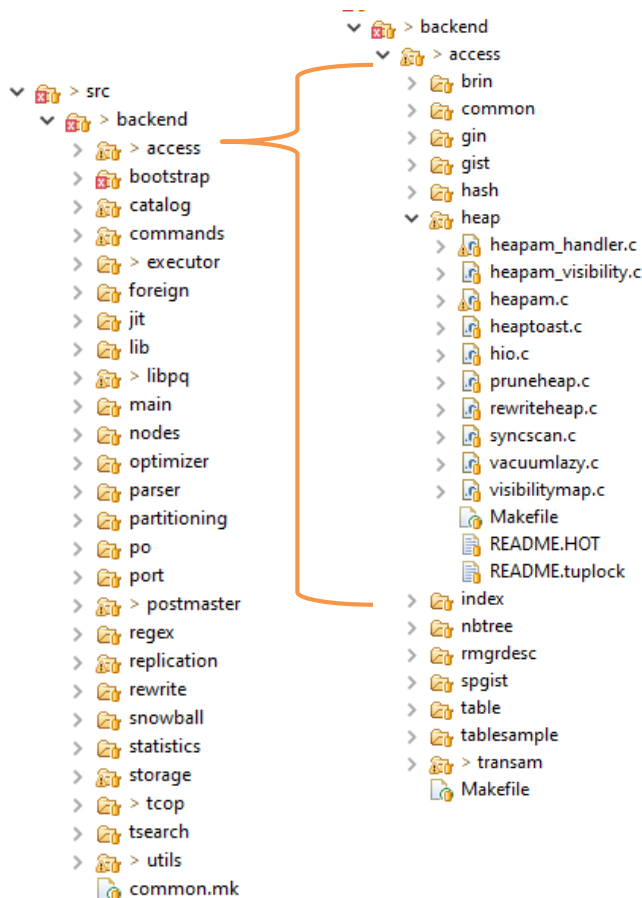
## Adding A New Backend Module in PostgreSQL



## Adding A New Module

### How To Add A New Module

- Most likely, the KMS module belongs to the backend of the PostgreSQL system, so it will reside in the “backend” folder.
- But there are many sub-folders under the backend folder and each sub-folder contains more folders.
- Where should you place the new module?
- There is no right or wrong answers.
- Put your new module in a place where you feel makes the most sense
- Once you decided on the place, copy the Makefile from another component that is in the same level as yours.
- Change the names to your module.





## Connect Your New Module

### Bootstrap vs Initialization

- Most of the PG module will have to go through either a bootstrap or initialization process or both.
- Bootstrap process happens during database initialization, (initdb). This means the database server is not really running at this stage.
  - This is only run **once** during initdb.
  - For example, the XLOG module, (the module responsible for WAL logs), has to go through a bootstrap process to generate the initial WAL files.
- Initialization process happens when a database is started (pg\_ctl start). After initialization, the database server remains running.
  - This is run **every time** the database is started.
  - For example, the XLOG module, is also required to go through an initialization process to register some space from shared memory.

### Consideration:

Does your new KMS module requires bootstrap or initialization or BOTH?





# Bootstrap Process

## Bootstrap

- Bootstrap is controlled by the source file “bootstrap.c”
- Bootstrap process starts from the “**AuxiliaryProcessMain()**” function in src/backend/bootstrap/bootstrap.c
- This may be a good location to add your own bootstrap logic if you need one.
- You simply create a new public function in your module and make sure it gets called during bootstrap process here

Src/backend/bootstrap/bootstrap.c

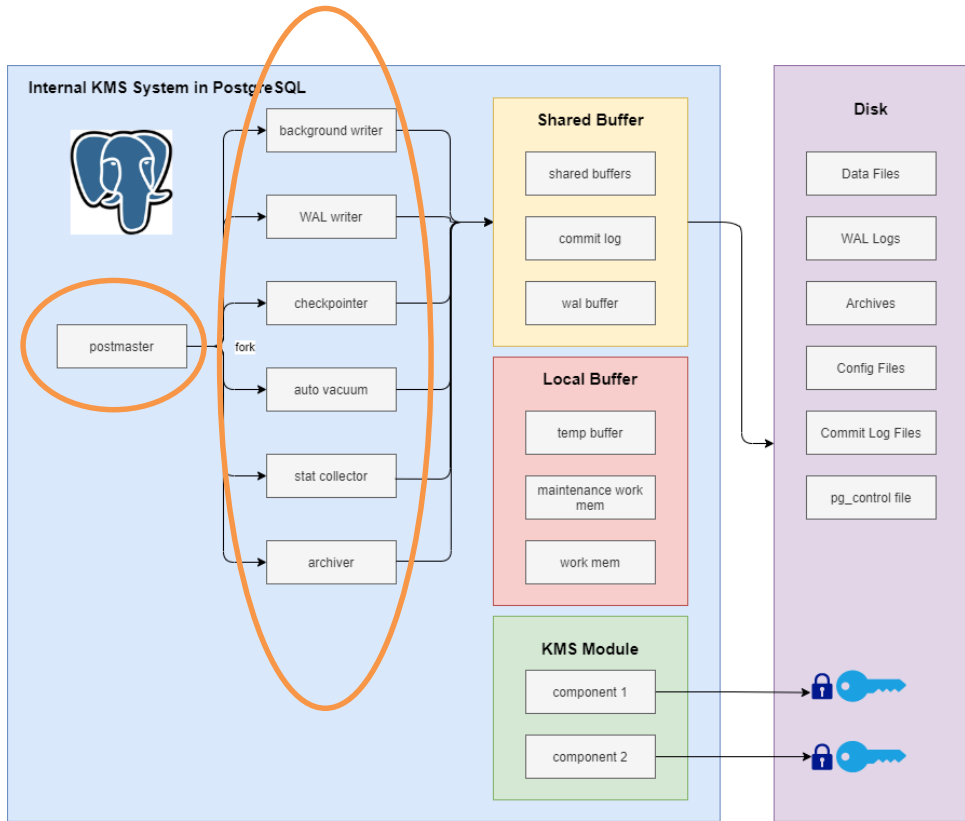
```
197 void
198 AuxiliaryProcessMain(int argc, char *argv[])
199 {
200     char    *progname = argv[0];
201     int      flag;
202     char    *userOption = NULL;
203
204     /*
205      * Initialize process environment (already done if under postmaster, but
206      * not if standalone).
207      */
208     if (!IsUnderPostmaster)
209         InitStandaloneProcess(argv[0]);
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
```



# Initialization Process

## Initialization

- A module's initialization normally happens at 2 places.
- When postmaster starts up
- When a standalone backend starts up





# Initialization Process

## Postmaster Initialization

- Remember, postmaster is the very first process that gets started.
- It will initialize almost all the backend components during startup.
- It all starts from the “**PostmasterMain()**” function in `src/backend/postmaster/postmaster.c`
- Here may be a good location to add your own initialization logics if you need one.

```
574 /*
575  * Postmaster main entry point
576  */
577 void
578 PostmasterMain(int argc, char *argv[])
579 {
580     int         opt;
581     int         status;
582     char        *userOption = NULL;
583     bool        listen_addr_saved = false;
584     int         i;
585     char        *output_config_variable = NULL;
586
587     InitProcessGlobals();
588
589     PostmasterPid = MyProcPid;
590
591     IsPostmasterEnvironment = true;
592
593     /*
594      * Remove old temporary files. At this point there can be no other
595      * PostgreSQL processes running in this directory, so this should be safe.
596      */
597     RemovePgTempFiles();
598
599     /*
600      * Initialize stats collection subsystem (this does NOT start the
601      * collector process!)
602      */
603     pgstat_init();
604
605     /*
606      * Initialize the autovacuum subsystem (again, no process start yet)
607      */
608     autovac_init();
609
610     /*
611      * Load configuration files for client authentication.
612      */
613     if (!load_hba())
614     {
615         /*
616          * It makes no sense to continue if we fail to load the HBA file,
617          * since there is no way to connect to the database in this case.
618          */
619         ereport(FATAL,
620                (errmsg("could not load pg_hba.conf")));
621     }
622 }
```



# Initialization Process

## Standalone Backend Initialization

- This is simply a process that is started by the postmaster as a standalone backend and it could be intended to work as one of the backend processes such as “bgwriter”, “checkpointer”, “Stat collector”..etc
- Although they are referred by different names, fundamentally they are all coming from one single binary file called “postgres”
- It all starts from the “**PostgresMain()**” function in `src/backend/tcop/postgres.c`
- Here may be a good location to add your own initialization logics if you need one. Make sure you do it after the **BaseInit()** as **BaseInit()** initializes file access, storage manager and buffer pool.

*Src/backend/tcop/postgres.c*

```
3777 void
3778 PostgresMain(int argc, char *argv[],
3779              const char *dbname,
3780              const char *username)
3781 {
3782     int firstchar;
3783     StringInfoData input_message;
3784     sigjmp_buf local_sigjmp_buf;
3785     volatile bool send_ready_for_query = true;
3786     bool disable_idle_in_transaction_timeout = false;
3787
3788     /* Initialize startup process environment if necessary. */
3789     if (!IsUnderPostmaster)
3790         InitStandaloneProcess(argv[0]);
3791
3792     SetProcessingMode(InitProcessing);
3793
3794
3795     /* Early initialization */
3796     BaseInit();
3797
3798     /*
3799      * Create a per-backend PGPROC struct in shared memory, except in the
3800      * EXEC_BACKEND case where this was done in SubPostmasterMain. We must do
3801      * this before we can use LWLocks (and in the EXEC_BACKEND case we already
3802      * had to do some stuff with LWLocks).
3803      */
3804     #ifdef EXEC_BACKEND
3805     if (!IsUnderPostmaster)
3806         InitProcess();
3807     #else
3808     InitProcess();
3809     #endif
3810
3811     /* We need to allow SIGINT, etc during the initial transaction */
3812     PG_SETMASK(&UnBlockSig);
3813
3814     /*
```



# Shared Memory Initialization

## Register Your Own Shared Memory From PG

- We know that PostgreSQL is a **multiple-process system**, meaning that there are many backend processes running behind the scene responsible for doing different things.
- They rely on **shared memory** to communicate between themselves to ensure the resource is accessed without conflict.
- Most likely, you will need to have something in the shared memory.

### Step 1: Define your struct

This is the structure of data that you would like to share among all backend processes. You can define it within your own module's header file (.h). Note that it is best to use finite data type so we know the size. Avoid using pointers here.

```
typedef struct ESLocalPage
{
    long    max_page_num;
    long    used_page_num;
    long    idle_page;
    LWLock  lock;
} ESLocalPage;
```

### Step 2: Declare a global variable using your struct

Make sure to declare it in global scale, not within a function.

**For Example:**

```
static ESLocalPage *es_local_pages = NULL;
```



# Shared Memory Initialization

## Register Your Own Shared Memory From PG

**Step 3:** Create a function to return the structure's size.  
Remember to also add this function's prototype in your header file so other components can access it.

**For example:**

```
Size es_page_size(void)
{
    return MAXALIGN(sizeof(ESLocalPage));
}
```

This is the global variable  
you declared in [step 2](#)

**Step 4:** Create a function to initialize your structure  
This is the function we will use to call PG's shared memory routine ([ShmemInitStruct](#)) to register shared memory for our component. You can give it any name in the first parameter of the function and give the size function in [step 2](#) in second parameter.

**For example:**

```
void es_page_init(void)
{
    bool    found;

    es_local_pages = (ESLocalPage *)
        ShmemInitStruct("ESlocal page", es_page_size(), &found);

    if(!found && es_local_pages)
    {
        memset(es_local_pages, 0, es_page_size());
        es_local_pages->max_page_num = 0L;
        es_local_pages->used_page_num = 0L;
        es_local_pages->idle_page = -1L;
    }
}
```



# Shared Memory Initialization

## Register Your Own Shared Memory From PG

So far, we have 1 structure and 2 function defined in our example. Next, we need to make them called from somewhere.

- es\_page\_size()
- es\_page\_init()

**Step 5:** Go to `src/backend/storage/ipc/ipci.c`, function `CreateSharedMemoryAndSemaphores()`

This is the function where PG tries to estimate the total size of shared memory block needed. You can see all other components are requesting share memory here:

Add a new line at the bottom of them after “`AsyncShmemSize()`”, but replacing the second parameter of “`add_size`” with the size function created in [step 3](#).  
**For example:**

`Size = add_size(size, es_page_size());`

```
95 void
96 CreateSharedMemoryAndSemaphores(void)
97 {
98     PGShmemHeader *shim = NULL;
```

*Src/backend/storage/ipc/ipci.c*

```
119     size = 100000;
120     size = add_size(size, PGSemaphoreShmemSize(numSemas));
121     size = add_size(size, SpinlockSemaSize());
122     size = add_size(size, hash_estimate_size(SHMEM_INDEX_SIZE,
123                                             sizeof(ShmemIndexEnt)));
124     size = add_size(size, BufferShmemSize());
125     size = add_size(size, LockShmemSize());
126     size = add_size(size, PredicateLockShmemSize());
127     size = add_size(size, ProcGlobalShmemSize());
128     size = add_size(size, XLOGShmemSize());
129     size = add_size(size, CLOGShmemSize());
130     size = add_size(size, CommitTsShmemSize());
131     size = add_size(size, SUBTRANSShmemSize());
132     size = add_size(size, TwoPhaseShmemSize());
133     size = add_size(size, BackgroundWorkerShmemSize());
134     size = add_size(size, MultiXactShmemSize());
135     size = add_size(size, LWLockShmemSize());
136     size = add_size(size, ProcArrayShmemSize());
137     size = add_size(size, BackendStatusShmemSize());
138     size = add_size(size, SInvalShmemSize());
139     size = add_size(size, FMSignalShmemSize());
140     size = add_size(size, ProcSignalShmemSize());
141     size = add_size(size, CheckpointerShmemSize());
142     size = add_size(size, AutoVacuumShmemSize());
143     size = add_size(size, ReplicationSlotsShmemSize());
144     size = add_size(size, ReplicationOriginShmemSize());
145     size = add_size(size, WalSndShmemSize());
146     size = add_size(size, WalRcvShmemSize());
147     size = add_size(size, ApplyLauncherShmemSize());
148     size = add_size(size, SnapMgrShmemSize());
149     size = add_size(size, BTreeShmemSize());
150     size = add_size(size, SyncScanShmemSize());
151     size = add_size(size, AsyncShmemSize());
```





# Shared Memory Initialization

## Register Your Own Shared Memory From PG

**Step 6:** In the same function, add your init function

Scroll down a little further within the same function in step 5, you can see a bunch of init functions for other components

Add a new line at the bottom of them again after AsyncShmemInit(), but with your own init function created in [step 4](#).

**For example:**

es\_page\_init();

That Is It!

```
95 void
96 CreateSharedMemoryAndSemaphores(void) Src/backend/storage/ipc/ipci.c
97 {
98     PGShmemHeader *shim = NULL;

232     InitPredicateLocks();
233
234     /*
235      * Set up process table
236      */
237     if (!IsUnderPostmaster)
238         InitProcGlobal();
239     CreateSharedProcArray();
240     CreateSharedBackendStatus();
241     TwoPhaseShmemInit();
242     BackgroundWorkerShmemInit();
243
244     /*
245      * Set up shared-inval messaging
246      */
247     CreateSharedInvalidationState();
248
249     /*
250      * Set up interprocess signaling mechanisms
251      */
252     PMSignalShmemInit();
253     ProcSignalShmemInit();
254     CheckpointerShmemInit();
255     AutoVacuumShmemInit();
256     ReplicationSlotsShmemInit();
257     ReplicationOriginShmemInit();
258     WalSndShmemInit();
259     WalRcvShmemInit();
260     ApplyLauncherShmemInit();
261
262     /*
263      * Set up other modules that need some shared memory space
264      */
265     SnapMgrInit();
266     BTreeShmemInit();
267     SyncScanShmemInit();
268     AsyncShmemInit();
```



## Add A New Config Parameter

### Add Your Own Params In Postgresql.conf

- You may also want to define your own configuration parameters for your new component.
- Configuration parameter is also referred as **GUC** (Global Unified Configuration).
- Follow these steps to add your own parameter

#### Step 1: Declare Your Global Config Variable

This is the global variable that you have to declare first within your component and initialize it to a default value. You can declare it within your component's source file.

**For example:**

```
bool        sharedsm_standby = false;
```



# Add A New Config Parameter

## Add Your Own Params In Postgresql.conf

### Step 2: Go to src/backend/utils/misc/guc.c

Depending on the data type of your parameters:

- Bool
- Int
- Real
- String
- Enum

You will need to locate the corresponding structure that defines your parameter.

For example, we want our parameter to be treated as a number so we will go to “`ConfigureNamesInt[]`” and add our parameter there:

*Src/backend/utils/misc/guc.c*

```
static struct config_int ConfigureNamesInt[] =
{
    {
        {"archive_timeout", PGC_SIGHUP, WAL_ARCHIVING,
         gettext_noop("Forces a switch to the next WAL file if a "
                     "new file has not been started within N seconds."),
         NULL,
         GUC_UNIT_S
        },
        &XLogArchiveTimeout,
        0, 0, INT_MAX / 2,
        NULL, NULL, NULL
    },

    /* Initialize mode for sharedsm */
    {
        {"sharedsm_standby", PGC_POSTMASTER, REPLICATION_STANDBY,
         gettext_noop("Set the mode for sharedsm, default off."),
         NULL
        },
        &sharedsm_standby,
        0, 0, 1,
        NULL, NULL, NULL
    },

    /* End-of-list marker */
    {
        {NULL, 0, 0, NULL, NULL}, NULL, 0, 0, 0, NULL, NULL, NULL
    }
};
```



# Add A New Config Parameter

## Add Your Own Params In Postgresql.conf

### Step 3: Categorize Your Parameter

Make sure you select the right category for your parameter using these 2 enum values:

- **PGC\_POSTMASTER** means the value can only be set when postmaster starts from config file or the command line. Read about other possible options in [guc.h](#)
- **REPLICATION\_STANDY** is the category of the parameter. The example belongs in replication category. Yours may not be. You can put “UNGROUPED” if appropriate. Read about all possible values in [guc\\_tables.h](#)

*Src/backend/utils/misc/guc.c*

```
/* Initialize mode for sharedsm */
{
    {"sharedsm_standby", PGC_POSTMASTER, REPLICATION_STANDBY,
     gettext_noop("Set the mode for sharedsm, default off."),
     NULL
    },
    &sharedsm_standby,
    0, 0, 1,
    NULL, NULL, NULL
},

/* End-of-list marker */
{
    {NULL, 0, 0, NULL, NULL}, NULL, 0, 0, 0, NULL, NULL, NULL
}
};
```



# Add A New Config Parameter

## Add Your Own Params In Postgresql.conf

### Step 4: Add the new parameter in the default postgresql.conf.sample

- Now that you have added a new GUC in the PostgreSQL system, the next thing you want to do is add your parameter in the default `postgresql.conf.sample`, so people know about this parameter.
- When you finish `initdb`, this sample conf file will be used as default in the new cluster
- This file is located in `src/backend/utils/postgresql.conf.sample`
- Add your parameter to the end of the file
- The hash (#) before the parameter name means the default value (0) will be used. You need to remove this hash, if you want to change the value to something else.

### Src/backend/utils/postgresql.conf.sample

```
#-----  
# CUSTOMIZED OPTIONS  
#-----  
  
# Add settings for extensions here  
  
#-----  
# SharedSM OPTIONS  
#-----  
#sharedsm_standby = 0                                # primary: 0, standby: 1
```

That Is It!



# The SQL Function

## What Is a SQL Function?

- A SQL function is ... a function that you can execute on your psql client terminal.
- PostgreSQL has a list of built-in SQL functions that you can execute using the **SELECT** command.
- In fact, we have used some of these SQL functions in previous lectures... To get the actual data file on disk, remember?
- For your new KMS module, you may also need to define one or more SQL functions to allow user to interact with the module.

For example, to trigger KMS module to do a key rotation.

## Examples of SQL Functions:

```
postgres=# select current_date;  
current_date  
-----  
2021-04-16  
(1 row)
```

```
postgres=# select pg_relation_filepath('test2');  
pg_relation_filepath  
-----  
base/12709/16387  
(1 row)
```



# Add Your Own SQL Function

## How To Add One For Your Module?

### Step 1: Identify an Unused OID Value

- Remember PostgreSQL references objects using something called a OID value?
- It is basically a number that uniquely identifies an object such as SQL function or tables within PostgreSQL.
- We need to know an OID value that is not currently being used by PG, so we can assign that to our new SQL function.
- To find out, run the script located in `src/include/catalog/unused_oids`
- And pick a good number to use

For example: 5566

```
caryh@HGPC01:~/highgo/git/projectv/sharedsm$ src/include/catalog/unused_oids
4 - 9
210
270 - 273
357
380 - 381
421
560 - 583
606
702 - 704
760 - 763
784 - 789
811 - 816
1177
1179 - 1180
1382 - 1383
1986 - 1987
2023
2030
2121
2137
2228
3432
3434 - 3435
3998
4035
4142
4187 - 4199
4225 - 4299
4388 - 4399
4532 - 4565
4572 - 4999
5022 - 5027
5032 - 5554
5556 - 5999
6015 - 6099
6103
6105
6107 - 6109
6116
6122 - 9999
```





# Add Your Own SQL Function

## How To Add One For Your Module?

### Step 2: Register your SQL function in pg\_proc.dat

- This step is to tell PG catalog that what kind of SQL function you are adding to PG.
- What is the input and output
- The name of the function
- The OID

#### For example:

a SQL function called `pg_rotate_cluster_passphrase` that takes no argument and returns a Boolean can be registered as:

```
# function for key managements
{ oid => '5566', descr => 'rotate cluster passphrase',
  proname => 'pg_rotate_cluster_passphrase',
  provolatile => 'v', prorettype => 'bool',
  proargtypes => '', prosrc => 'pg_rotate_cluster_passphrase' },
```

#### Src/include/catalog/pg\_proc.dat

```
#-----
#
# pg_proc.dat
#   Initial contents of the pg_proc system catalog.
#
# Portions Copyright (c) 1996-2020, PostgreSQL Global Development Group
# Portions Copyright (c) 1994, Regents of the University of California
#
# src/include/catalog/pg_proc.dat|
#
#-----
```

Simply add this  
block here at the  
end of the file

Visit the official documentation on pg\_proc for more information:  
<https://www.postgresql.org/docs/current/catalog-pg-proc.html>



# Add Your Own SQL Function

## How To Add One For Your Module?

### Step 3: Add your SQL function prototype declaration

- Like a regular C function, a SQL function also needs a prototype declaration so that other modules are aware of this function.
- You can add your prototype in the function manager module in “src/include/utils/fmgrprotos.h”
- A SQL function always has an abstract return type of “datum” and an abstract input argument of “PG\_FUNCTION\_ARGS”

For example:

```
extern Datum pg_rotate_cluster_passphrase(PG_FUNCTION_ARGS);
```

```
* fmgrprotos.h
```

Src/include/utils/fmgrprotos.h

```
#ifndef FMGRPROTOS_H
#define FMGRPROTOS_H

#include "fmgr.h"

extern Datum heap_tableam_handler(PG_FUNCTION_ARGS);
extern Datum byteaout(PG_FUNCTION_ARGS);
extern Datum charout(PG_FUNCTION_ARGS);
extern Datum namein(PG_FUNCTION_ARGS);
extern Datum nameout(PG_FUNCTION_ARGS);
extern Datum int2in(PG_FUNCTION_ARGS);
extern Datum int2out(PG_FUNCTION_ARGS);
extern Datum int2vectorin(PG_FUNCTION_ARGS);
extern Datum int2vectorout(PG_FUNCTION_ARGS);
extern Datum int4in(PG_FUNCTION_ARGS);
extern Datum int4out(PG_FUNCTION_ARGS);
extern Datum regprocin(PG_FUNCTION_ARGS);
extern Datum regprocout(PG_FUNCTION_ARGS);
extern Datum textin(PG_FUNCTION_ARGS);
extern Datum textout(PG_FUNCTION_ARGS);
extern Datum tidin(PG_FUNCTION_ARGS);
extern Datum tidout(PG_FUNCTION_ARGS);
```

Add your prototype  
at the end of this file



# Add Your Own SQL Function

## How To Add One For Your Module?

### Step 4: Implement The SQL Function

- Now, you just need to implement the SQL function and fill it with logics
- This function should be implemented within your module somewhere
- Remember, this example function does not take any input, so we are not doing anything with `PG_FUNCTION_ARGS`.
- This function returns a Boolean as we have registered to catalog.
- So, in order to return properly, we need to wrap it with a macro, `PG_RETURN_BOOL(true)`. This macro basically makes “datum” to represent a Boolean value as we require.

```
/*
 * SQL function to rotate the cluster passphrase. This function assumes that
 * the cluster_passphrase_command is already reloaded to the new value.
 * All internal keys are wrapped by the new passphrase and saved to the disk.
 * To update all crypto keys atomically we save the newly wrapped keys to the
 * temporary directory, pg_cryptokeys_tmp, and remove the original directory,
 * pg_cryptokeys, and rename it. These operation is performed without the help
 * of WAL. In the case of failure during rotationpg_cryptokeys directory and
 * pg_cryptokeys_tmp directory can be left in incomplete status. We recover
 * the incomplete situation by checkIncompleteRotation.
 */
Datum
pg_rotate_cluster_passphrase(PG_FUNCTION_ARGS)
{
    PgAeadCtx *ctx;
    CryptoKey newkeys[KMGR_MAX_INTERNAL_KEYS] = {0};
    char passphrase[KMGR_MAX_PASSPHRASE_LEN];
    uint8 new_kekenc[PG_AEAD_ENC_KEY_LEN];
    uint8 new_kekmac[PG_AEAD_MAC_KEY_LEN];
    int passlen;

    /* Rename to the original directory */
    if (rename(KMGR_TMP_DIR, KMGR_DIR) != 0)
        ereport(ERROR,
                (errmsg("could not rename directory \"%s\" to \"%s\": %m",
                        KMGR_TMP_DIR, KMGR_DIR)));
    fsync_fname(KMGR_DIR, true);

    LWLockRelease(KmgrFileLock);

    pg_free_aead_ctx(ctx);
    PG_RETURN_BOOL(true);
}
```



## More On Datum and PG\_FUNCTION\_ARGS

### What are these?

- Datum and PG\_FUNCTION\_ARGS are just abstractions to all of the data types.
- So instead of worrying about what datatype to use for input or output when designing a PG function, you can just give it either datum or PG\_FUNCTION\_ARGS
- When you implement the function, you then cast them to what you need.
- Taking “pg\_relation\_size” SQL function as example:
  - First input argument is cast to OID
  - Second input argument is cast to Text
  - And it returns a 64-bit integer value

```
/* Datum
pg_relation_size(PG_FUNCTION_ARGS)
{
    Oid          relOid = PG_GETARG_OID(0);
    text         *forkName = PG_GETARG_TEXT_PP(1);
    Relation     rel;
    int64        size;

    rel = try_relation_open(relOid, AccessShareLock);

    /*
     * Before 9.2, we used to throw an error if the relation didn't exist, but
     * that makes queries like "SELECT pg_relation_size(oid) FROM pg_class"
     * less robust, because while we scan pg_class with an MVCC snapshot,
     * someone else might drop the table. It's better to return NULL for
     * already-dropped tables than to throw an error and abort the whole query.
     */
    if (rel == NULL)
        PG_RETURN_NULL();

    size = calculate_relation_size(&(rel->rd_node), rel->rd_backend,
                                  forkname_to_number(text_to_cstring(forkName)));

    relation_close(rel, AccessShareLock);

    PG_RETURN_INT64(size);
}
```



```
        #deselection at the end - add back the deselected mirror modifier object
        mirror_ob.select= 1
        modifier_ob.select=1
        bpy.context.scene.objects.active = modifier_ob
        print("Selected" + str(modifier_ob)) # modifier ob is the active ob
        #mirror_ob.select = 0
        $one = bpy.context.selected_objects[0]
        #bpy.data.objects[one.name].select = 1
    except:
        print("please select exactly two objects, the last one gets the modifier unless its not a mod")
```

```
----- OPERATOR CLASSES -----
Mirror Tool
```

```
class MirrorX(bpy.types.Operator):
    """This adds an X mirror to the selected object"""
    bl_idname = "object.mirror_mirror_x"
    bl_label = "Mirror X"
```

```
    classmethod
    def poll(cls, context):
        return context.active_object is not None
```



## Internal KMS Consideration



## Define A Scope

### What Features To Support?

- We have learned that KMS normally involves in 6 steps in a key's life cycle.
  - Key Generation (required)
  - Key Storage (required)
  - Key Renewal
  - Key Rotation (required)
  - Key Revocation
  - Key Destruction
- I would say 3 out of the above 6 are required as minimum



Key  
Generation

Key  
Storage

Key  
Renewal

Key  
Rotation

Key  
Revocation

Key  
Destruction

瀚高  
HIGH GO



融知与行 瀚且高远

THANKS